



# TensorFlow

Курс “Практическое применение по TensorFlow”

Шигапова Фирюза Зинатуллаевна

1-й семестр, 2019 г.



<https://github.com/Firyuza/TensorFlowPractice>

# Override Backward-Propagation

- Model contains **non-differentiable** operations
- You must define own custom flow in backprop
- You must define own custom forward flow

# Override Backward-Propagation. Model

```
stddev = 1e-1
labels = [0, 1]
data = [1, 2]

labels_ph = tf.placeholder(shape=(2), dtype=tf.int32)
data_ph = tf.placeholder(shape=(2), dtype=tf.float32)

c = tf.Variable(tf.truncated_normal(shape=[2], stddev=stddev))
k = tf.Variable(tf.truncated_normal(shape=[2], stddev=stddev))

d = tf.add(data_ph, c)
e = tf.add(c, k)
a = tf.multiply(d, e)

loss = tf.losses.softmax_cross_entropy(logits=a, onehot_labels=labels_ph)
opt = tf.train.AdamOptimizer(learning_rate=0.001, beta1=0.9, beta2=0.999, epsilon=0.1)
```

# Override Backward-Propagation. Model

```
stddev = 1e-1
labels = [0, 1]
data = [1, 2]

labels_ph = tf.placeholder(shape=(2), dtype=tf.int32)
data_ph = tf.placeholder(shape=(2), dtype=tf.float32)

c = tf.Variable(tf.truncated_normal(shape=[2], stddev=stddev))
k = tf.Variable(tf.truncated_normal(shape=[2], stddev=stddev))


d = tf.add(data_ph, c)
e = tf.add(c, k)
a = tf.multiply(d, e)

loss = tf.losses.softmax_cross_entropy(logits=a, onehot_labels=labels_ph)
opt = tf.train.AdamOptimizer(learning_rate=0.001, beta1=0.9, beta2=0.999, epsilon=0.1)
```

Let's Override backprop for this operation

# Override Backward-Propagation

Write function that will replace `tf.add` operation:



# Override Backward-Propagation

Say TensorFlow for which operation you want to override Backward-Propagation:

Use **gradient\_override\_map** that belongs to Graph

```
with graph.gradient_override_map({"PyFunc": rnd_name}):  
    return tf.py_func(forward_func, inputs, [np.float32], name=name)
```

It is **tf.add** operation

**All variables** that take part in **tf.add** operation

Type of returned data from forward func

# Override Backward-Propagation

```
def forward_func(c, k, d):  
    e = np.add(c, k)  
  
    return e.astype(np.float32)
```

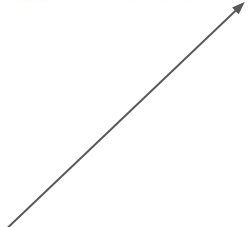
Need to pass all variables that take part NOT only in forward but in backprop too



Compute gradients

# Override Backward-Propagation

```
tf.RegisterGradient(rnd_name)(backprop_func)
```




Define gradient function for  
an operation



Function that computes gradients in  
backprop flow

# Graph



## Backprop: chain rule

$$\frac{\partial L}{\partial c} = \frac{\partial d}{\partial c} \cdot \frac{\partial L}{\partial d} + \frac{\partial e}{\partial c} \cdot \frac{\partial L}{\partial e} = e \cdot \frac{\partial L}{\partial a} + d \cdot \frac{\partial L}{\partial a} = (e + d) \cdot \frac{\partial L}{\partial a},$$

where  $\frac{\partial L}{\partial d} = \frac{\partial a}{\partial d} \cdot \frac{\partial L}{\partial a} = e \cdot \frac{\partial L}{\partial a}$  and  $\frac{\partial L}{\partial e} = \frac{\partial a}{\partial e} \cdot \frac{\partial L}{\partial a} = d \cdot \frac{\partial L}{\partial a}$

$$\frac{\partial L}{\partial k} = \frac{\partial e}{\partial k} \cdot \frac{\partial L}{\partial e} = d \cdot \frac{\partial L}{\partial a}$$

Note: the gradient  $\frac{\partial L}{\partial a}$  is already calculated, we get it in backprop, in my example this variable called "grad" in `backprop_func`.

# Override Backward-Propagation

```
def backprop_func(op, grad):  
    c = op.inputs[0]  
    k = op.inputs[1]  
    d = op.inputs[2]  
    e = tf.add(c, k)  
    return (e + d) * grad, d * grad, e * grad
```

Compute partial derivative for each passed variable

- **op** contains all passed variables;
- **grad** — the flown gradient from the back propagation.

# Override Backward-Propagation

```
def custom_add(c, k, d, graph, name=None):  
    with tf.name_scope(name, "AddGrad", [c, k, d]) as name:  
        # Need to generate a unique name to avoid duplicates  
        # if you have more than one custom gradients:  
        rnd_name = 'PyFuncGrad' + str(np.random.randint(0, 1e+8))  
  
        inputs = [c, k, d]  
  
        tf.RegisterGradient(rnd_name)(backprop_func)  
        with graph.gradient_override_map({"PyFunc": rnd_name}):  
            return tf.py_func(forward_func, inputs, [np.float32], name=name)
```

# Override Backward-Propagation

Explicitly call **compute\_gradients** and apply to optimizer!!!

```
grads = opt.compute_gradients(loss, tf.trainable_variables())  
grads = list(grads)  
train_op = opt.apply_gradients(grads_and_vars=grads)
```

# Override Backward-Propagation

```
stddev = 1e-1
labels = [0, 1]
data = [1, 2]

labels_ph = tf.placeholder(shape=(2), dtype=tf.int32)
data_ph = tf.placeholder(shape=(2), dtype=tf.float32)

c = tf.Variable(tf.truncated_normal(shape=[2], stddev=stddev))
k = tf.Variable(tf.truncated_normal(shape=[2], stddev=stddev))

d = tf.add(data_ph, c)
e = custom_add(c, k, d, tf.get_default_graph()) # tf.add(c, k)
a = tf.multiply(d, e)

loss = tf.losses.softmax_cross_entropy(logits=a, onehot_labels=labels_ph)
opt = tf.train.AdamOptimizer(learning_rate=0.001, beta1=0.9, beta2=0.999, epsilon=0.1)

grads = opt.compute_gradients(loss, tf.trainable_variables())
grads = list(grads)
train_op = opt.apply_gradients(grads_and_vars=grads)
```

# tf.stop\_gradient

Stop to flow backprop flow further for the given operation

```
tf.stop_gradient(  
    input,  
    name=None  
)
```



# Clothing Retrieval with Visual Attention Model

Zhonghao Wang <sup>1</sup>, Yujun Gu <sup>1</sup>, Ya Zhang <sup>1✉</sup>, Jun Zhou <sup>2</sup>, Xiao Gu <sup>2</sup>

<sup>1</sup> *Cooperative Medianet Innovation Center, Shanghai Jiao Tong University, Shanghai, China*

<sup>2</sup> *Institute of Image Communication and Network Engineering, Shanghai Jiao Tong University, Shanghai, China*

*{kindredcain,yjgu,ya\_zhang,zhoujun,gugu97}@sjtu.edu.cn*

They describe attention network that generates **Bernoulli** series has to be multiplied with another feature map.

Bernoulli is **not differentiable** => custom backpropagation

<https://arxiv.org/pdf/1710.11446.pdf>

# Control dependencies

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1...m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

# Control dependencies

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots x_m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i$$

// mini-batch mean

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2$$

// mini-batch variance

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$$

// normalize

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i)$$

// scale and shift


**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

These ops **will not** be evaluated during one train step

# Control dependencies

```
assign_mean = self.mean.assign(mean)
assign_variance = self.variance.assign(variance)
```


- Need to say evaluate it **First** and then train step.
- These ops are not explicitly called during train step evaluation.
- These nodes are independent in Graph.



**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

# Control dependencies

contains these UPDATE operations:  
assign\_mean, assign\_variance



```
extra_update_ops = tf.get_collection(tf.GraphKeys.UPDATE_OPS)
with tf.control_dependencies(extra_update_ops):
    train_op = optimizer.minimize(loss)
```

OR

```
extra_update_ops = tf.get_collection(tf.GraphKeys.UPDATE_OPS)
session.run([train_op, extra_update_ops], ...)
```

# Control dependencies. What will be printed?

Only zeros.

Operation is ***not*** created within Control Dependency context!

```
x = tf.Variable(0.0)
x_plus_1 = tf.assign_add(x, 1)

with tf.control_dependencies([x_plus_1]):
    y = x
init = tf.initialize_all_variables()

with tf.Session() as session:
    init.run()
    for i in range(5):
        print(session.run(y))
```

# Control dependencies. Solution

```
x = tf.Variable(0.0)
x_plus_1 = tf.assign_add(x, 1)

with tf.control_dependencies([x_plus_1]):
    y = tf.identity(x)
init = tf.initialize_all_variables()
```

OR

```
with tf.Session() as session:
    init.run()
    for i in range(5):
        print(session.run(y))
```

```
x = tf.Variable(0.0)
x_plus_1 = tf.assign_add(x, 1)


with tf.control_dependencies([x_plus_1]):
    y = x_plus_1
init = tf.initialize_all_variables()
```

```
with tf.Session() as session:
    init.run()
    for i in range(5):
        print(session.run(y))
```



# Operations by themselves

For example Center Loss: “A Discriminative Feature Learning Approach for Deep Face Recognition”



**Fig. 2.** The distribution of deeply learned features in (a) training set (b) testing set, both under the supervision of softmax loss, where we use 50K/10K train/test splits. The points with different colors denote features from different classes. **Best viewed in color.** (Color figure online)



# Operations by themselves

For example, Center Loss: “A Discriminative Feature Learning Approach for Deep Face Recognition”

$$\mathcal{L}_C = \frac{1}{2} \sum_{i=1}^m \|\mathbf{x}_i - \mathbf{c}_{y_i}\|_2^2$$

The  $\mathbf{c}_{y_i} \in \mathbb{R}^d$  denotes the  $y_i$ th class center of deep features.

$$\Delta \mathbf{c}_j = \frac{\sum_{i=1}^m \delta(y_i = j) \cdot (\mathbf{c}_j - \mathbf{x}_i)}{1 + \sum_{i=1}^m \delta(y_i = j)} \quad \mathbf{c}_j^{t+1} = \mathbf{c}_j^t - \alpha \cdot \Delta \mathbf{c}_j^t$$

Centers update is **independent** operation.

For updating centers values need to **explicitly** evaluate this op!

# Operations by themselves


For example Center Loss: “A Discriminative Feature Learning Approach for Deep Face Recognition”

```
centers = tf.get_variable(name_centers, [nrof_classes, nrof_features], dtype=tf.float32,  
    initializer=tf.constant_initializer(0), trainable=False)  
  
self.update_centers_op = tf.scatter_sub(centers, label, diff)  
  
self.sess.run(  
    [self.train_op, self.update_centers_op],  
    feed_dict=feed_dict)
```

← Evaluate **update\_centers\_op** with **train\_op**

# Train several networks independently

For example, Several GAN: “SGAN: An Alternative Training of Generative Adversarial Networks”.



# Train several networks independently

1. Create Graph for each network separately.

Train several networks that are defined in the **same** Graph is **impossible**.

2. Copy values of one network to another:

```
vars_1 = [var for var in train_vars1
           if '[needed scope name]' in var.name]

with self.session_2.as_default():
    vars_2 = [var for var in train_vars2
               if '[needed scope name]' in var.name]
    for var_idx, var in enumerate(vars_1):
        vars_2[var_idx].load(
            vars_1[var_idx].eval(session=self.session_1),
            self.session_2)
```

# Train several networks independently

- assign
- `tf.Variable`
- `tf.contrib.copy_graph.copy_variable_to_graph`

They make **reference**, not a **deep copy**.


For making a deep copy of trainable variables into another one use **load**.

# TensorFlow Queue

1. Always control that **queue is completely filled**.
2. Always be sure that **validation data is not** used as **training data** and vice versa.
3. During **dequeue** in **multithreading** way care about transformation: all random transformation has to be **RANDOM**.

# tf.gather

Take data from the given params according to indices.



# tf.gather

If code is running on **GPU**: indices can be out of range, **zeros** will be returned.

```
a = tf.constant([1, 2, 3, 4, 5])  
b = tf.gather(a, [0, 1, 6]) # [1 2 0]
```

If code is running on **CPU**: indices **cannot** be out of range, otherwise throws **error**.

```
a = tf.constant([1, 2, 3, 4, 5])  
b = tf.gather(a, [0, 1, 6])  
# InvalidArgumentError: indices[2] = 6 is not in [0, 5) [Op:GatherV2]
```