



TensorFlow

Курс “Практическое применение по TensorFlow”

Шигапова Фирюза Зинатуллаевна

1-й семестр, 2019 г.



<https://github.com/Firyuza/TensorFlowPractice>

TensorFlow 2.0 installation

1. Install Anaconda <https://docs.anaconda.com/anaconda/install/#>



The screenshot shows the Anaconda documentation website. On the left is a navigation sidebar with a tree structure. The main content area on the right is titled "Installing on macOS" and contains a list of steps for installation, including downloading the macOS installer and clicking the Install button. At the bottom, there is a small image of the "Install Anaconda3" window.

Navigation sidebar:

- Home
- Anaconda Enterprise 5
- Anaconda Enterprise 4
- Anaconda Distribution
 - Installation
 - Installing on Windows
 - Installing on macOS
 - Installing on Linux
 - Installing on Linux POWER
 - Installing in silent mode
 - Verifying your installation
 - Anaconda installer file hashes
 - Updating from older versions
 - Uninstalling Anaconda

Main content area:

Installing on macOS

You can install Anaconda using either the graphical installer ("wi; unsure, choose the graphical install.

macOS graphical install

1. Download the graphical [macOS installer](#) for your version of Py
2. OPTIONAL: [Verify data integrity with MD5 or SHA-256](#). For mc [verification?](#).
3. Double-click the downloaded file and click continue to start th
4. Answer the prompts on the Introduction, Read Me, and Licens
5. Click the Install button to install Anaconda in your home user c

Install Anaconda3 window:

Standard Install on "Macintosh HD"

This will take 2.13 GB of space on your computer.

Click Install to perform a standard installation of this software

TensorFlow 2.0 installation

2. Create conda environment:

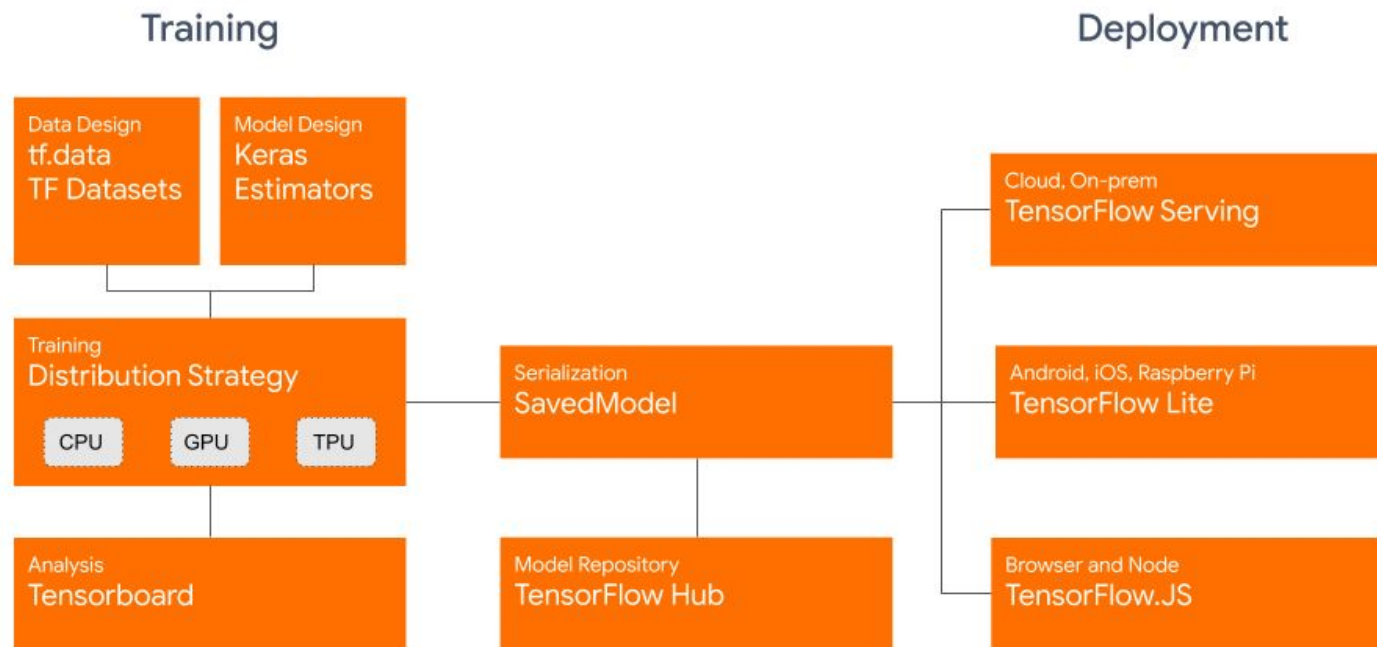
```
conda create -n tensorflow_2 python=3.x
```

3. Install TensorFlow into created environment (tensorflow 2.0):

```
source activate tensorflow_2
```

```
pip install tensorflow or pip install tensorflow-gpu (Now by default tensorflow 2.0)
```

<https://www.tensorflow.org/install/pip>



TensorFlow 1.x vs TensorFlow 2.0

TensorFlow 1.x	TensorFlow 2.0
	Eager execution
Static Graph	tf.function, AutoGraph
Session	Eager execution
Queue runner	tf.data
tf.train.Coordinator	tf.distribute.Strategy
tf.control_dependencies	<i>Does not exist anymore</i>
Slim, Keras, Estimator	Keras, Estimator
	tf.GradientTape
	TFX, TF Lite (and micro), TFjs,

Why TensorFlow 1.x is alive?

While eager execution makes development and debugging more interactive, TensorFlow graph execution has advantages for distributed training, performance optimizations, and production deployment. However, writing graph code can feel different than writing regular Python code and more difficult to debug.

TensorFlow 1.x vs. TensorFlow 2.0

```
in_a = tf.placeholder(dtype=tf.float32, shape=(2))
in_b = tf.placeholder(dtype=tf.float32, shape=(2))
```

TensorFlow 1.x

```
def forward(x):
    with tf.variable_scope("matmul", reuse=tf.AUTO_REUSE):
        W = tf.get_variable("W", initializer=tf.ones(shape=(2,2)),
                             regularizer=tf.contrib.layers.l2_regularizer(0.04))
        b = tf.get_variable("b", initializer=tf.zeros(shape=(2)))
        return W * x + b
```

```
out_a = forward(in_a)
out_b = forward(in_b)
```

```
reg_loss = tf.losses.get_regularization_loss(scope="matmul")
```

```
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    outs = sess.run([out_a, out_b, reg_loss],
                     feed_dict={in_a: [1, 0], in_b: [0, 1]})
```

TensorFlow 2.x

```
W = tf.Variable(tf.ones(shape=(2,2)), name="W")
b = tf.Variable(tf.zeros(shape=(2)), name="b")
```

```
def forward(x):
    return W * x + b
```

```
out_a = forward([1,0])
print(out_a)
```


tf.function and AutoGraph

- **tf.function** allows to transform Python syntax into high-performance TensorFlow Static Graph
- Or “defines a **computations** as a **graph** of TensorFlow operations”
- **tf.function** “accelerates” part of code but doesn’t allow to debug at runtime
- Using **tf.function** **AutoGraph** automatically applied (tf.autograph)

tf.function and AutoGraph

```
v = tf.Variable(1.0)
@tf.function
def f():
    v.assign(2.0)
    return v.read_value()
```

```
print(f()) # Always prints 2.0.
```

```
def foo(x):
    if x > 0:
        y = x * x
    else:
        y = -x
    return y
```

```
converted_foo = tf.autograph.to_graph(foo)
```

```
x = tf.constant(1)
y = converted_foo(x) # converted_foo is a TensorFlow Op-like.
assert tf.is_tensor(y)
```

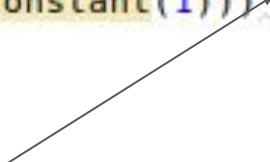


@tf.function uses it



tf.function and AutoGraph

```
@tf.function
def f(x):
    if x > 0:
        x = x + 1
    return x

tf.config.experimental_run_functions_eagerly(True)
print(f(tf.constant(1)))
```



To debug @tf.function

Build Model

1. Functional API

```
self.conv3x3_a = tf.keras.layers.Conv2D(filters, (kernel_size, kernel_size), padding='same')
self.conv3x3_a_gn = GroupNorm(filters)
self.conv3x3_b = tf.keras.layers.Conv2D(filters, (kernel_size, kernel_size), padding='same')
self.conv3x3_b_gn = GroupNorm(filters)
self.max_pool = tf.keras.layers.MaxPool2D(pool_size=(max_pool_size, max_pool_size))
```

2. Sequential API

```
self.sequential_model = tf.keras.Sequential()
self.sequential_model.add(tf.keras.layers.Conv2D(filters, (kernel_size, kernel_size), padding='same'))
self.sequential_model.add(GroupNorm(filters))
self.sequential_model.add(tf.keras.layers.Conv2D(filters, (kernel_size, kernel_size), padding='same'))
self.sequential_model.add(GroupNorm(filters))
self.sequential_model.add(tf.keras.layers.MaxPool2D(pool_size=(max_pool_size, max_pool_size)))
```

Run model via Keras API

It has the **same** API as it is used in **TensorFlow 1.x**

Next slide 

Keras. Compile model

Define necessary components for model **training**:

- optimizer
- loss function
- metrics

```
# Configure a model for categorical classification.  
model.compile(optimizer=tf.keras.optimizers.RMSprop(0.01),  
              loss=tf.keras.losses.CategoricalCrossentropy(),  
              metrics=[tf.keras.metrics.CategoricalAccuracy()])
```

Keras. Fit model

Define dataset for training and validation phases:

- training data (**tf.data** can be used)
- validation data (**tf.data** can be used)
- epochs
- batch size

```
model.fit(data, labels, epochs=10, batch_size=32,  
          validation_data=(val_data, val_labels))
```

Keras. Evaluate model

Returns loss and metrics values on valid/test set


```
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
```


Keras. Predict model


Returns predictions of trained model at inference

```
predictions = model.predict(test_images)
```

Custom “components”


- Layers
 - Models
 - Loss
- 
- Have **similar** creating logic
- Callbacks
 - Accuracy

Build complex model via Keras



Build complex model via Keras

Create custom Layer



Create custom Layer

tf.keras.layers.Layer

1. Inherit *custom* layer from **tf.keras.layers.Layer**

```
class UNetDownBlock(tf.keras.layers.Layer):  
    def __init__(self, filters, kernel_size, max_pool_size):  
        super(UNetDownBlock, self).__init__()
```

2. Create **trainable** variables in **constructor** or in **build** method

```
class UNetDownBlock(tf.keras.layers.Layer):  
    def __init__(self, filters, kernel_size, max_pool_size):  
        super(UNetDownBlock, self).__init__()   
        self.conv3x3_a = tf.keras.layers.Conv2D(filters, (kernel_size, kernel_size), padding='same')  
        self.conv3x3_a_gn = GroupNorm(filters)  
        self.conv3x3_b = tf.keras.layers.Conv2D(filters, (kernel_size, kernel_size), padding='same')  
        self.conv3x3_b_gn = GroupNorm(filters)  
        self.max_pool = tf.keras.layers.MaxPool2D(pool_size=(max_pool_size, max_pool_size))  
  
    def call(self, input_tensor, is_training=True):  
        output = self.conv3x3_a(input_tensor)  
        output = self.conv3x3_a_gn(output)  
        output_1 = self.conv3x3_b(output)  
        output_1 = self.conv3x3_b_gn(output_1)  
        output = self.max_pool_1(output_1)  
  
        return output_1, output
```

3. In **call** method write **forward** operations

tf.keras.layers.Layer

1. Use **Layers** as outputs for **model** creating

```
model = tf.keras.Model(inputs=inputs, outputs=outputs)
```

2. Use **standard** keras API:

- compile
- fit
- evaluate
- etc.

Different architectures can be wrapped as single keras Layer

- ResNet
- Variational AutoEncoder
- etc.

tf.keras.models.Model

1. Wrap custom Layers into **Model** (tf.keras.models.Model)
2. Build **tf.keras.Model** as tf.keras.layers.Layer:
 - in **constructor** or in **build** method create object of custom **Layer**
 - in **call** method write **forward** operations of the **Model**
3. In **train** pipeline:
 - create Model object
 - build object

```
self.unet = UNet()
self.unet.build((None, cfg.train.image_size,  cfg.train.image_size, 3))
```

4. At train step just use created Model object for calling **forward operations**

```
for image, label in dataset:
    with tf.GradientTape() as tape:
        segmentation_map = self.unet(image, True)
```


tf.keras.models.Model

5. Or use **standard** keras API:

- compile
- fit
- etc.

Custom Loss

1. Inherit *loss* from `tf.keras.losses.Loss`

```
class MaxMargin(tf.keras.losses.Loss):  
    def __init__(self, delta, from_logits=False,  
                  reduction=tf.keras.losses.Reduction.AUTO,  
                  name='max_margin'):  
        super(MaxMargin, self).__init__(reduction=reduction,  
                                         name=name)
```

2. Create needed **variables** in **constructor** or in **build** method

3. In **call** method write **forward** operations

```
def call(self, v1, v2):  
    norm = tf.sqrt(tf.reduce_sum(tf.square(tf.subtract(v1, v2, axis=1), axis=1), axis=1))  
  
    return tf.maximum(0., self.delta - norm)
```

Custom Loss

4. In train pipeline create Loss object

```
self.max_margin_loss = MaxMargin(cfg.train.delta)
```

5. At train step just use created Loss object

```
total_loss += self.max_margin_loss(emb1, emb2)
```

Custom Accuracy

1. Inherit from **tf.keras.metrics.Metric**
2. In constructor add variable for tracking custom accuracy value

```
class BinaryTruePositives(tf.keras.metrics.Metric):  
    def __init__(self, name='binary_true_positives', **kwargs):  
        super(BinaryTruePositives, self).__init__(name=name, **kwargs)  
        self.true_positives = self.add_weight(name='tp', initializer='zeros')
```

3. Update accuracy value using **update_state** method

```
def update_state(self, y_true, y_pred, sample_weight=None):  
    y_true = tf.cast(y_true, tf.bool)  
    y_pred = tf.cast(y_pred, tf.bool)  
  
    values = tf.logical_and(tf.equal(y_true, True), tf.equal(y_pred, True))  
    values = tf.cast(values, self.dtype)  
  
    self.true_positives.assign_add(tf.reduce_sum(values))
```

Custom Accuracy

4. Override result method for getting current accuracy value

```
def result(self):  
    return self.true_positives
```

5. Work in pipeline as with pre-made keras accuracies

Custom Callback

1. **Inherit** custom callback from **tf.keras.callbacks.Callback**
2. **Override** needed methods:
 - on_train_batch_begin
 - on_train_batch_end
 - on_epoch_begin
 - etc

Custom Callback

```
class MyCustomCallback(tf.keras.callbacks.Callback):

    def on_train_batch_begin(self, batch, logs=None):
        print('Training: batch {} begins at {}'.format(batch, datetime.datetime.now().time()))

    def on_train_batch_end(self, batch, logs=None):
        print('Training: batch {} ends at {}'.format(batch, datetime.datetime.now().time()))

    def on_test_batch_begin(self, batch, logs=None):
        print('Evaluating: batch {} begins at {}'.format(batch, datetime.datetime.now().time()))

    def on_test_batch_end(self, batch, logs=None):
        print('Evaluating: batch {} ends at {}'.format(batch, datetime.datetime.now().time()))
```

tf.GradientTape

- Track operations for automatic differentiation
- All **trainable** variable are **automatically** watched
- For manually watching use method **watch()**

tf.GradientTape. How To Use?

- Create context using `tf.GradientTape()`

```
with tf.GradientTape() as tape:  
    segmentation_map = self.unet(image, True)
```


- Operations recorded by *GradientTape* if they executed **within** this **context** manager or by calling `watch()`
- Compute **gradient** and **apply** them to **optimizer**

```
grads = tape.gradient(loss_value, self.unet.trainable_variables)  
self.optimizer.apply_gradients(zip(grads, self.unet.trainable_variables))
```

tf.GradientTape. How To Use?

```
with tf.GradientTape() as tape:  
    segmentation_map = self.unet(image, True)  
  
    loss_value = self.unet_loss(label, segmentation_map)  
  
grads = tape.gradient(loss_value, self.unet.trainable_variables)  
self.optimizer.apply_gradients(zip(grads, self.unet.trainable_variables))
```

Immediately compute
gradients



Try toy example with **x = tf.Variable(5.)**

tf.GradientTape. “Watching”

```
x = tf.constant(3.0)
with tf.GradientTape() as g:
    g.watch(x)
    y = x * x
dy_dx = g.gradient(y, x) # 6.0
```

Non-trainable variable => have to manually “watch”

tf.GradientTape. Lifecycle

- By default **tf.GradientTape()** allows to call gradient method once and then tape is removed
- For multiple calls:

```
x = tf.constant(3.0)
with tf.GradientTape(persistent=True) as g:
    g.watch(x)
    y = x * x
    z = y * y
dz_dx = g.gradient(z, x) # 108.0
dy_dx = g.gradient(y, x) # 6.0
```

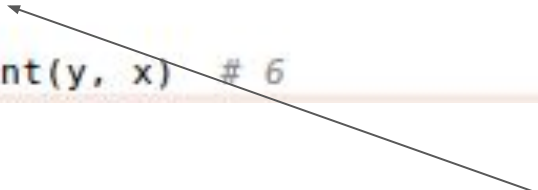
tf.GradientTape. Control gradients

```
x = tf.Variable(3.0)
with tf.GradientTape(watch_accessed_variables=False) as g:
    y = x * x
dy_dx = g.gradient(y, x) # None
```



No gradients will be computed since watching is **disabled**

```
x = tf.Variable(3.0)
with tf.GradientTape(watch_accessed_variables=False) as g:
    g.watch(x)
    y = x * x
dy_dx = g.gradient(y, x) # 6
```



Manually watch needed operations